

Name: Himanehu Kumar
Roll no.: 2022215

Assignment-4

Question 1

a) Forward pass of Convolutional Layer

1) Given

image dimension: $M \times N$ with P channels

kernel size: $K \times K$

stride $S = 1$

No padding

\therefore The dimension of the resulting feature map $O_h \times O_w$ are given by

$$O_h = M - K + 1 \quad O_w = N - K + 1$$

Thus output feature map dimension are:

$$(M - K + 1) \times (N - K + 1)$$

2) For each output pixel:

\rightarrow $K \times K$ height for each of the P input channels, requiring $K^2 P$ multiplication.

\rightarrow The result are summed, requiring $K^2 P - 1$ addition.

$$\therefore \text{total operation} = K^2 P + K^2 P - 1$$

$$= 2K^2 P - 1 \Rightarrow 2 \cdot K^2 P - 1$$

For large K , this simplified to

$$O(K^2 \cdot P)$$

3) with Q kernels, the total no. of operation is

$$\text{Total operation} = Q \cdot (M - K + 1)(N - K + 1) \cdot (2K^2P - 1)$$

Big-O Complexity:

→ for General Case:

$$O(Q \cdot M \cdot N \cdot K^2 \cdot P)$$

→ Assuming $\min(M, N) \gg K$:

$$O(Q \cdot M \cdot N \cdot K^2 \cdot P)$$

b) 1) Assignment Step

→ Compute the distance of each data point to all k Centroids

→ Assign the data points to the nearest Centroid.

2) Update Step

→ Compute the mean of all points in each cluster.

→ Update the cluster centroid positions.

3) Determine the Optimal no. of Cluster.
We will use Elbow method.

→ Compute the inertia for different value of k .

→ Plot inertia vs k

→ optimal k : The elbow point where the decrease in inertia slow significantly.

4) Random initialization may lead to poor convergence:

→ k -mean minimize a non-convex objective function

→ May converge to local minimum

using k -mean++ for initialization improve the likelihood of better results by ensuring centroids are well-separated initially.

ML ASSIGNMENT-4

Question2

a)


```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # Initial data points (X) and centroids
5 X = np.array([
6     [5.1, 3.5], [4.9, 3.0], [5.8, 2.7], [6.0, 3.0], [6.7, 3.1], [4.5, 2.3], [6.1, 2.8],
7     [5.2, 3.2], [5.5, 2.6], [5.0, 2.0], [8.0, 0.5], [7.5, 0.8], [8.1, -0.1], [2.5, 3.5],
8     [1.0, 3.0], [4.5, -1.0], [3.0, -0.5], [5.1, -0.2], [6.0, -1.5], [3.5, -0.1], [4.0, 0.0],
9     [6.1, 0.5], [5.4, -0.5], [5.3, 0.3], [5.8, 0.6]
10 ])
11
12 # Initial centroids
13 initial_centroids = np.array([[3.0, 3.0], [2.0, 2.0]])
14 centroids = initial_centroids.copy()
15 k = centroids.shape[0]
16 max_iterations = 100
17 tolerance = 1e-4
18
19 def euclidean_distance(a, b):
20     return np.sqrt(np.sum((a - b) ** 2))
21
22 def assign_clusters(X, centroids):
23     clusters = []
24     for x in X:
25         distances = [euclidean_distance(x, centroid) for centroid in centroids]
26         closest_centroid = np.argmin(distances)
27         clusters.append(closest_centroid)
28     return np.array(clusters)
29
30 def update_centroids(X, clusters, k):
31     new_centroids = []
32     for i in range(k):
33         cluster_points = X[clusters == i]
34         if len(cluster_points) > 0:
35             new_centroid = np.mean(cluster_points, axis=0)
36         else:
37             new_centroid = centroids[i]
38         new_centroids.append(new_centroid)
39     return np.array(new_centroids)
40
41 def has_converged(old_centroids, centroids, tolerance):
42     return np.all(np.linalg.norm(centroids - old_centroids, axis=1) < tolerance)
43
```

```

1 # KMeans Algorithm
2 for iteration in range(max_iterations):
3     # Step 1: Assignment Step
4     clusters = assign_clusters(X, centroids)
5
6     # Step 2: Update Step
7     new_centroids = update_centroids(X, clusters, k)
8
9     # Step 3: Convergence Check
10    if has_converged(centroids, new_centroids, tolerance):
11        print(f"Converged after {iteration + 1} iterations.")
12        break
13
14    # Update centroids for the next iteration
15    centroids = new_centroids
16
17 # Plotting
18 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
19
20 # Plot initial clusters
21 ax1.scatter(X[initial_clusters == 0], X[initial_clusters == 0, 1], color='blue', label='Cluster 1')
22 ax1.scatter(X[initial_clusters == 1], X[initial_clusters == 1, 1], color='green', label='Cluster 2')
23 ax1.scatter(initial_centroids[:, 0], initial_centroids[:, 1], color='red', marker='x', s=100, label='Initial Centroids')
24 ax1.set_title('Initial Clusters and Centroids')
25 ax1.legend()
26
27 # Plot final clusters
28 ax2.scatter(X[clusters == 0], X[clusters == 0, 1], color='blue', label='Cluster 1')
29 ax2.scatter(X[clusters == 1], X[clusters == 1, 1], color='green', label='Cluster 2')
30 ax2.scatter(centroids[:, 0], centroids[:, 1], color='red', marker='x', s=100, label='Final Centroids')
31 ax2.set_title('Final Clusters and Centroids')
32 ax2.legend()
33
34 plt.show()
35
36 # Print final centroid values
37 print("Final centroids:\n", centroids)

```

b)



Converged after 3 iterations.

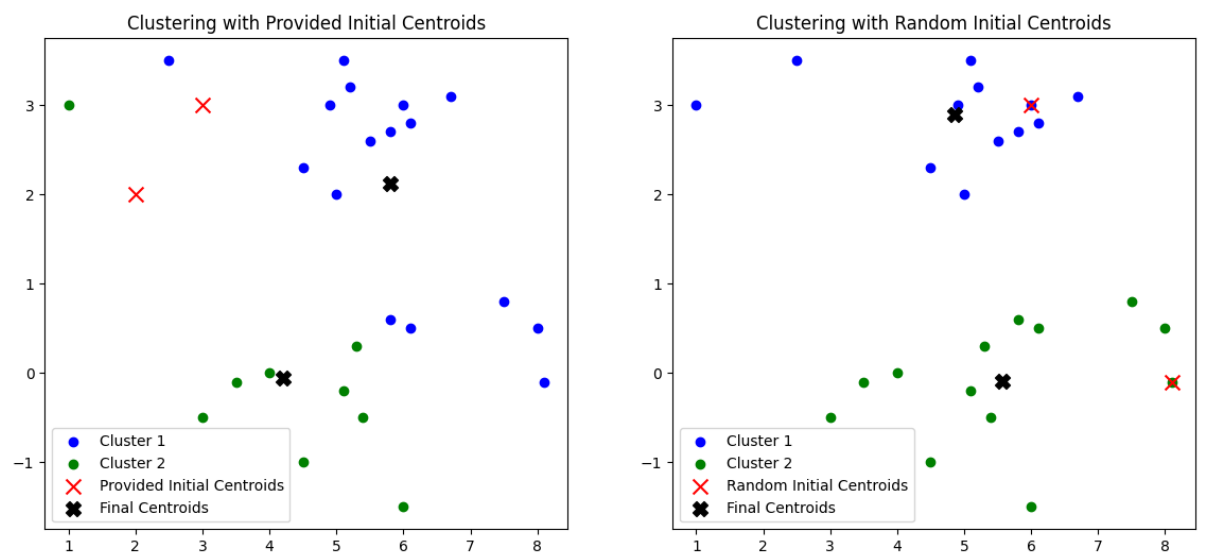
Final centroids:

```

[[ 5.8    2.125 ]
 [ 4.2   -0.05555556]]

```

c)




Final centroids with provided initial centroids:

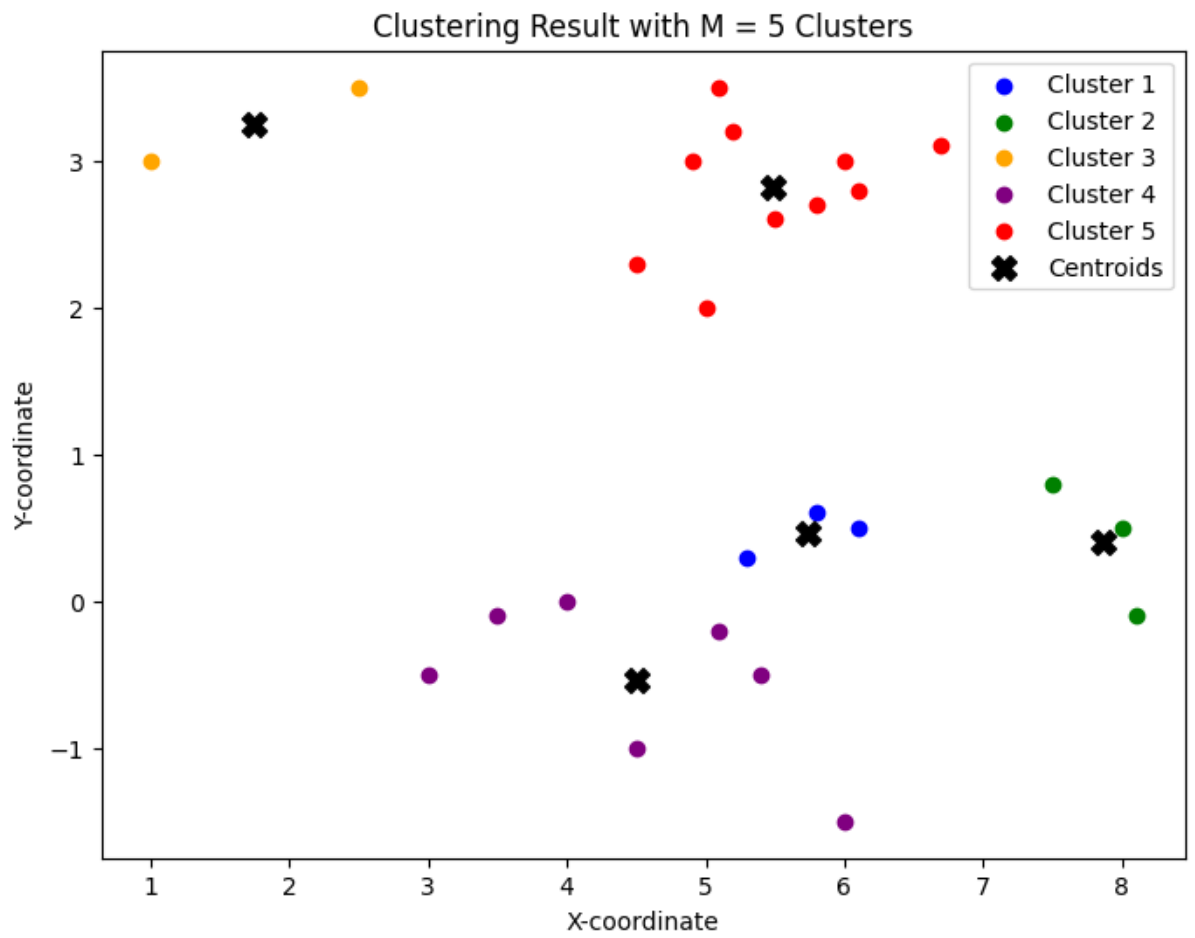
```
[[ 5.8    2.125 ]
 [ 4.2   -0.0555556]]
```

Final centroids with random initial centroids:


```
[[ 4.85833333 2.89166667]
 [ 5.56153846 -0.09230769]]
```

d)





When $m=5$



When $m=3$

Question-3

a)

```
1 import torch
2 from torchvision import datasets, transforms
3 from torch.utils.data import DataLoader, Subset
4 from sklearn.model_selection import train_test_split
5 import numpy as np
6
7 # Define transformations
8 transform = transforms.Compose([
9     transforms.ToTensor(),
10    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
11 ])
12
13 # Load CIFAR-10 dataset
14 dataset = datasets.CIFAR10(root='./data', train=True, transform=transform, download=True)
15
16 # Choose 3 classes (e.g., 0: Airplane, 1: Automobile, 2: Bird)
17 selected_classes = [0, 1, 2]
18 indices = [i for i, (_, label) in enumerate(dataset) if label in selected_classes]
19
20 # Split indices into train and test
21 train_indices, val_indices = train_test_split(indices, test_size=0.2, stratify=[dataset.targets[i] for i in indices])
22
23 # Define custom datasets
24 class CustomDataset(torch.utils.data.Dataset):
25     def __init__(self, dataset, indices):
26         self.dataset = dataset
27         self.indices = indices
28
29     def __len__(self):
30         return len(self.indices)
31
32     def __getitem__(self, idx):
33         return self.dataset[self.indices[idx]]
34
35 train_dataset = CustomDataset(dataset, train_indices)
36 val_dataset = CustomDataset(dataset, val_indices)
37
38 # Data loaders
39 train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True)
40 val_loader = DataLoader(val_dataset, batch_size=64, shuffle=False)
41
42 # Test dataset (subset from original test data)
43 test_dataset = datasets.CIFAR10(root='./data', train=False, transform=transform, download=True)
44 test_indices = [i for i, (_, label) in enumerate(test_dataset) if label in selected_classes]
45 test_dataset = CustomDataset(test_dataset, test_indices)
46 test_loader = DataLoader(test_dataset, batch_size=64, shuffle=False)
47
```

b) Training Dataset:


airplane



automobile



bird



Validation Dataset:



airplane



automobile



bird



c)

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4
5 class CNN(nn.Module):
6     def __init__(self):
7         super(CNN, self).__init__()
8         self.conv1 = nn.Conv2d(3, 16, kernel_size=5, stride=1, padding=1)
9         self.pool1 = nn.MaxPool2d(kernel_size=3, stride=2)
10        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=0)
11        self.pool2 = nn.MaxPool2d(kernel_size=3, stride=3)
12        self.fc1 = nn.Linear(512, 16)
13        self.fc2 = nn.Linear(16, 3)
14
15    def forward(self, x):
16        x = self.pool1(F.relu(self.conv1(x)))
17        x = self.pool2(F.relu(self.conv2(x)))
18        x = x.view(x.size(0), -1) # Flatten the output
19        # print(x.shape)
20        x = F.relu(self.fc1(x))
21        x = self.fc2(x)
22        return x
23
```

d)

Epoch 1/15: Train Loss: 0.7344, Train Accuracy: 67.66% | Val Loss: 0.5643, Val Accuracy: 78.13%
Epoch 2/15: Train Loss: 0.5168, Train Accuracy: 79.62% | Val Loss: 0.4633, Val Accuracy: 81.37%
Epoch 3/15: Train Loss: 0.4581, Train Accuracy: 82.04% | Val Loss: 0.4613, Val Accuracy: 81.87%
Epoch 4/15: Train Loss: 0.4175, Train Accuracy: 83.71% | Val Loss: 0.3972, Val Accuracy: 84.37%
Epoch 5/15: Train Loss: 0.3817, Train Accuracy: 84.91% | Val Loss: 0.4417, Val Accuracy: 82.47%
Epoch 6/15: Train Loss: 0.3639, Train Accuracy: 86.05% | Val Loss: 0.3733, Val Accuracy: 85.40%

Epoch 7/15: Train Loss: 0.3429, Train Accuracy: 86.58% | Val Loss: 0.3676, Val Accuracy: 85.80%

Epoch 8/15: Train Loss: 0.3270, Train Accuracy: 87.40% | Val Loss: 0.3698, Val Accuracy: 85.83%

Epoch 9/15: Train Loss: 0.3132, Train Accuracy: 87.82% | Val Loss: 0.3477, Val Accuracy: 86.53%

Epoch 10/15: Train Loss: 0.2995, Train Accuracy: 88.18% | Val Loss: 0.4132, Val Accuracy: 84.13%

Epoch 11/15: Train Loss: 0.2864, Train Accuracy: 88.95% | Val Loss: 0.3584, Val Accuracy: 86.20%

Epoch 12/15: Train Loss: 0.2754, Train Accuracy: 89.39% | Val Loss: 0.3417, Val Accuracy: 86.93%


Epoch 13/15: Train Loss: 0.2727, Train Accuracy: 89.40% | Val Loss: 0.3548, Val Accuracy: 86.73%

Epoch 14/15: Train Loss: 0.2659, Train Accuracy: 89.65% | Val Loss: 0.3487, Val Accuracy: 87.00%

Epoch 15/15: Train Loss: 0.2502, Train Accuracy: 90.08% | Val Loss: 0.3416, Val Accuracy: 86.93%

Model saved to: model\cnn_model.pth

e)



CNN Test Accuracy: 88.30%

CNN Test F1-Score: 0.8831



f)

Epoch 1/15: Train Loss: 0.6868, Train Accuracy: 72.10% | Val Loss: 0.6446, Val Accuracy: 74.83%

Epoch 2/15: Train Loss: 0.5729, Train Accuracy: 77.37% | Val Loss: 0.5826, Val Accuracy: 77.93%

Epoch 3/15: Train Loss: 0.5094, Train Accuracy: 80.24% | Val Loss: 0.5701, Val Accuracy: 78.43%

Epoch 4/15: Train Loss: 0.4681, Train Accuracy: 82.15% | Val Loss: 0.5591, Val Accuracy: 78.67%

Epoch 5/15: Train Loss: 0.4375, Train Accuracy: 83.33% | Val Loss: 0.5653, Val Accuracy: 77.90%

Epoch 6/15: Train Loss: 0.3960, Train Accuracy: 84.86% | Val Loss: 0.5512, Val Accuracy: 80.00%

Epoch 7/15: Train Loss: 0.3845, Train Accuracy: 85.56% | Val Loss: 0.5691, Val Accuracy: 79.03%

Epoch 8/15: Train Loss: 0.3494, Train Accuracy: 87.12% | Val Loss: 0.5435, Val Accuracy: 80.40%

Epoch 9/15: Train Loss: 0.3227, Train Accuracy: 88.11% | Val Loss: 0.5544, Val Accuracy: 80.00%

Epoch 10/15: Train Loss: 0.3099, Train Accuracy: 88.58% | Val Loss: 0.5859, Val Accuracy: 80.60%

Epoch 11/15: Train Loss: 0.2895, Train Accuracy: 89.31% | Val Loss: 0.6132, Val Accuracy: 79.33%

Epoch 12/15: Train Loss: 0.2609, Train Accuracy: 90.58% | Val Loss: 0.6765, Val Accuracy: 79.53%

Epoch 13/15: Train Loss: 0.2558, Train Accuracy: 90.95% | Val Loss: 0.6034, Val Accuracy: 80.50%

Epoch 14/15: Train Loss: 0.2234, Train Accuracy: 92.16% | Val Loss: 0.6340, Val Accuracy: 79.57%

Epoch 15/15: Train Loss: 0.2091, Train Accuracy: 92.69% | Val Loss: 0.7023, Val Accuracy: 77.67%

Model saved to: model\cnn_model.pth

Test Accuracy (MLP): 79.20%

Test F1-Score (MLP): 0.7927

Performance Comparison:

MLP Test Accuracy: 79.20%, MLP Test F1-Score: 0.7927




g)

Test Accuracy (MLP): 80.43%

Test F1-Score (MLP): 0.8050

Performance Comparison:

MLP Test Accuracy: 80.43%, MLP Test F1-Score: 0.8050



In test accuracy and the F1 score, the CNN model beats the MLP model, which suggests the fact that convolutional layers are essential for success with image data. The exclusive benefit of spatial relationship capture lies with the CNN-that is not the architecture of the MLP, which ignores them.

Nevertheless, MLP model was not bad in performance overall but is not comparable with CNN when image classification task is concerned. The performance differences again stress the need for different architectures relied upon different applications, which in this case is the CNN for image task-type applications.